

# **Linear Machine Learning**

Janis Keuper

# Introduction to ML



#### **Basic Types of Machine Learning Algorithms**

**Supervised Learning** 

**Unsupervised Learning** 

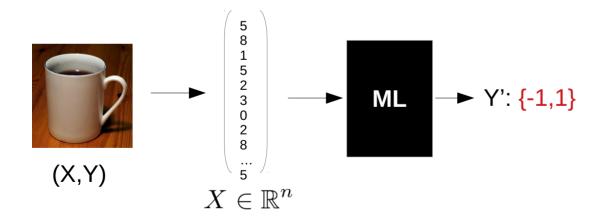
**Reinforcement Learning** 

- Labeled data
- Direct and quantitative evaluation
- Learn model from "ground truth" examples
- Predict unseen examples

# **Recall Classification**



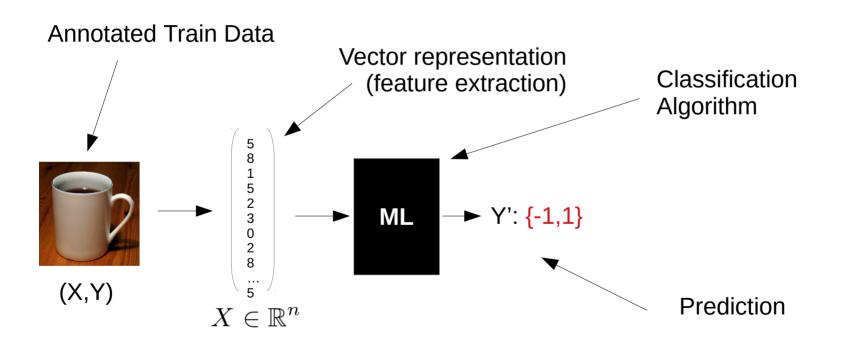
## **Supervised Learning: Annotated Training Data**



# **Recall Classification**



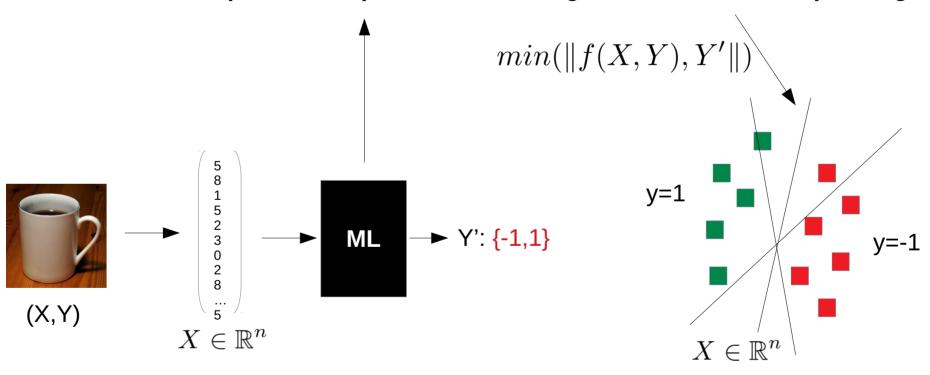
## **Supervised Learning: Annotated Training Data**



# **Recall Classification**



#### **LEARNING**: is a optimization problem → Finding the best function separating



# **Example: MNIST**

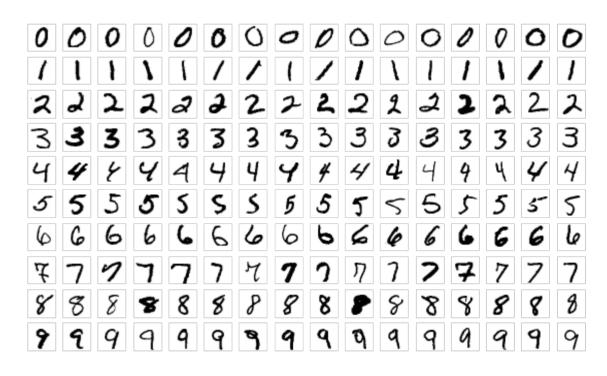


#### The MNIST hand written digits classification Problem

The MNIST database (Modified National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems.

#### **Data specs:**

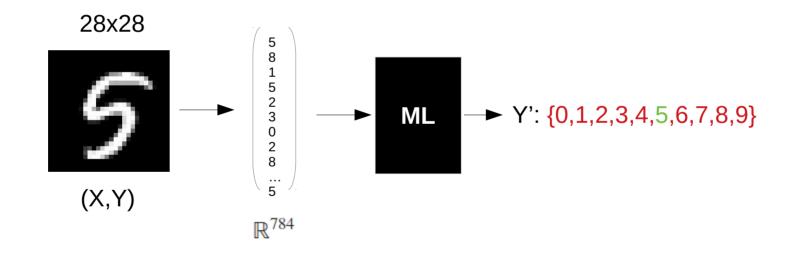
- · 10 Classes (digest 0-9)
- · 28x28 gray scale images
- · 60000 train samples
- · 10000 test samples



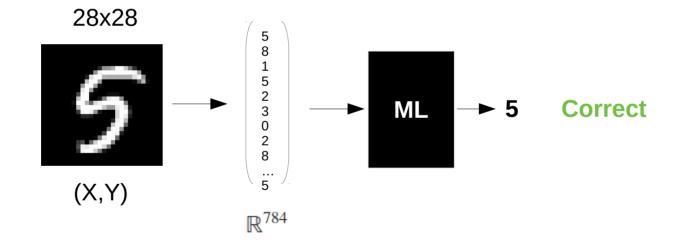
# **Example: MNIST**



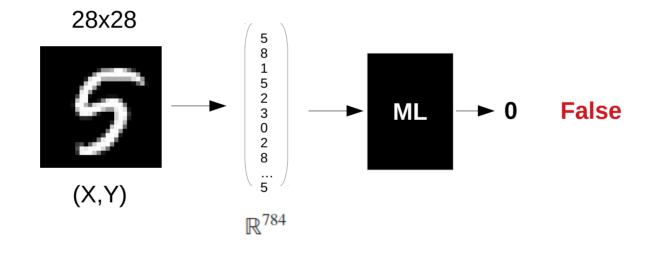
## The MNIST hand written digits classification Problem











# **Recall: Evaluation**



#### Basic evaluation of a model:

**Train error:** measure of how well the model predicts the given labels

$$Err_{train} := \frac{1}{|X_{train}|} \sum_{x_i \in X_{train}} |f(x_i) - y_i|$$

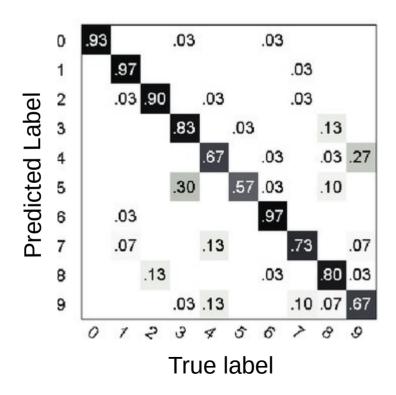
low train error is the necessary condition for a "good" model

Test error: same as train error: low test error is the sufficient condition

$$Err_{test} := \frac{1}{|X_{test}|} \sum_{x_i \in X_{test}} |f(x_i) - y_i|$$



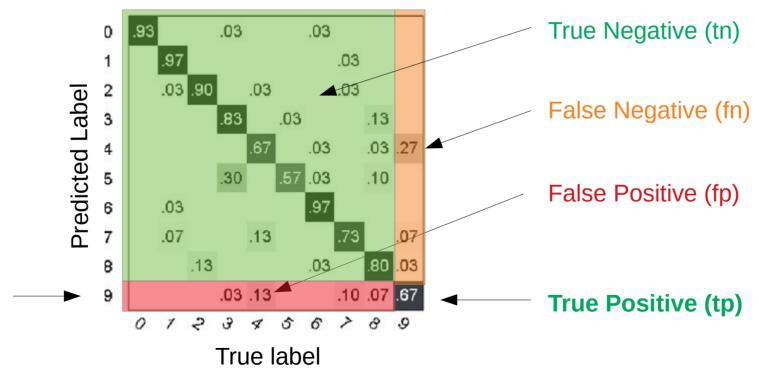
## **Confusion Matrix and True and False Positives/Negatives**





## **Confusion Matrix and True and False Positives/Negatives**

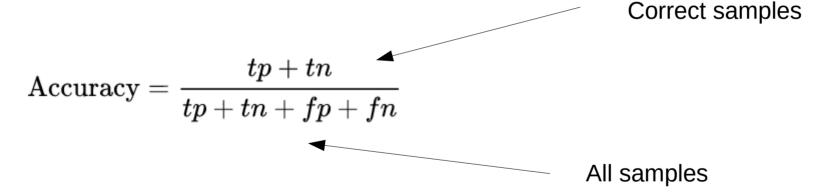
Example for true digit "9"





## **Accuracy:**

#### Most commonly used error metric:





#### **Problems with Accuracy Unbalanced classes:**

If the prior probability of one class is much higher than others, *fp* will have little impact.

$$ext{Accuracy} = rac{tp+tn}{tp+tn+fp+fn}$$

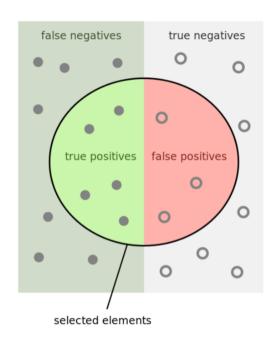
Extreme example: if 90% of the digits are "1", classifying every digit to "1" will will have 90% accuracy!



#### **Precision and Recall**

$$ext{Precision} = rac{tp}{tp+fp}$$

$$ext{Recall} = rac{tp}{tp+fn}$$



How many selected items are relevant?

How many relevant items are selected?

[image by wikipedia]



## F-Measure or balanced F-score is the harmonic mean of precision and recall:

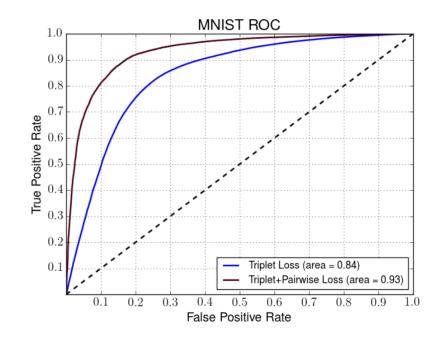
$$F = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



#### **Receiver Operating Characteristic Curve**

A receiver operating characteristic curve, or ROC curve, is a graphical plot that illustrates the diagnostic ability of a **binary classifier** system as its discrimination threshold is varied.

Example: comparing two different models

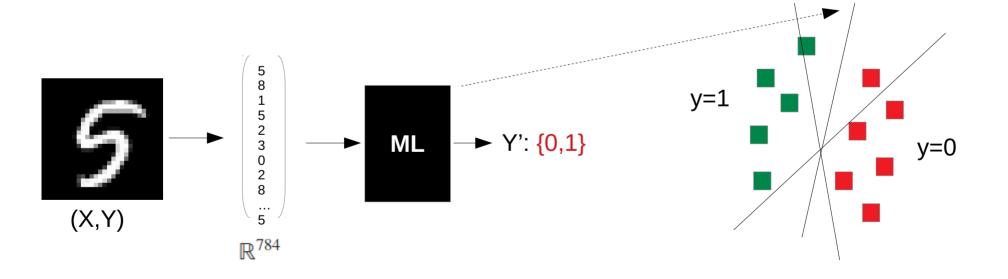




A Simple Linear Model: binary classification

**Example: "5" vs "other digit"** 

Model: hyper plane





A Simple Linear Model: binary classification

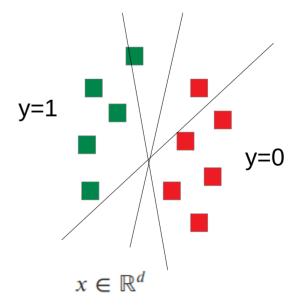
Parameterization of prediction function f with d-dimensional data as:

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$

With data samples  $x \in \mathbb{R}^d$ 

Model parameters  $w \in \mathbb{R}^d$ 

Model: hyper plane





A Simple Linear Model: binary classification

Parameterization of prediction function f with d-dimensional data as:

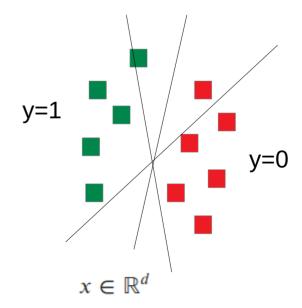
$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$

With data samples  $x \in \mathbb{R}^d$ 

Model parameters  $w \in \mathbb{R}^d$ 

How to find the parameters?

Model: hyper plane





#### **Optimization problem to find parameters**

$$\underset{w}{\arg\min} \sum_{i=0}^{N} L(y_i, w^T x_i)$$

#### With a differential Loss function like

$$L(y = 1, y') := \frac{1}{1 + e^{-y'}}$$

$$L(y = 0, y') := 1 - L(y = 1, y')$$



## **Optimization problem to find parameters**

$$\underset{w}{\arg\min} \sum_{i=0}^{N} L(y_i, w^T x_i)$$

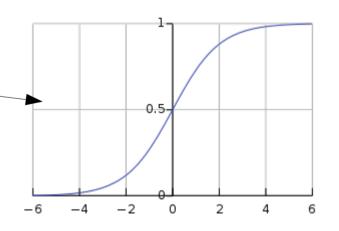
With a differential Loss function like: logistic function

$$L(y = 1, y') := \frac{1}{1 + e^{-y'}}$$

$$L(y = 0, y') := 1 - L(y = 1, y')$$

- pseudo probability: Out put always between 0 and 1
- Apply threshold function on probability that class label =1

Only one of many possible Loss functions, but common choice





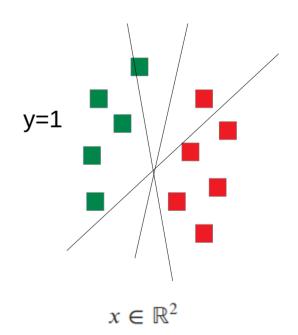
Goal: find w to minimize  $\arg \min \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

$$\underset{w}{\arg\min} \sum_{i=0}^{N} L(y_i, w^T x_i)$$



Goal: find w to minimize  $\underset{w}{\operatorname{arg \, min}} \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

2D Example:



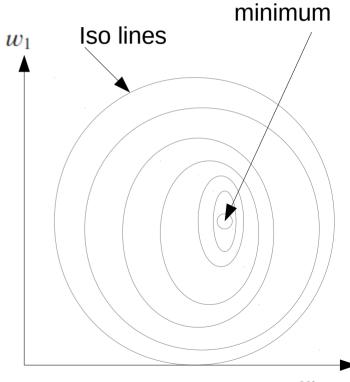
**Feature Space** 



Goal: find w to minimize  $\underset{w}{\operatorname{arg \, min}} \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

## 2D Example:

- How many (and which) parameters do we have to find?
- *L* spans a (loss) surface in the d-dimensional space of the data X (parameter space)
- We can evaluate L at each point w
- We can compute the gradient at each point w in L (assuming L to be Lipschitz)



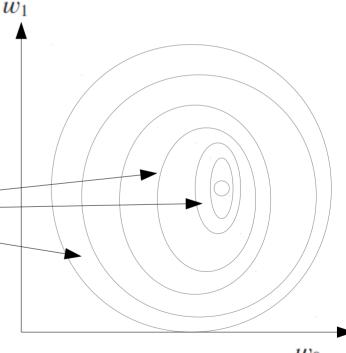
**Model Parameter Space** 



Goal: find w to minimize  $\arg \min \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

## 2D Example:

- How many (and which) parameters do we have to find?
- L spans a (loss) surface in the d-dimensional space of the data X (parameter space)
- We can evaluate L at each point w
- We can compute the gradient at each point w in L (assuming L to be Lipschitz)



 $w_0$ 



Goal: find w to minimize  $\underset{w}{\operatorname{arg min}} \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

## 2D Example:

- How many (and which) parameters do we have to find?
- *L* spans a (loss) surface in the d-dimensional space of the data X (parameter space)
- We can evaluate L at each point w
- We can compute the gradient at each point w in L (assuming L to be Lipschitz)  $\nabla L = \frac{dL}{dw}$

 $\nabla L$ 

 $w_1$ 

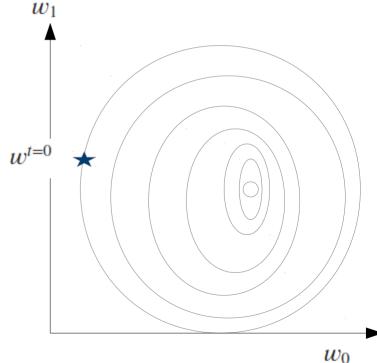
 $w_0$ 



Goal: find w to minimize  $\arg \min \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

#### **Gradient Descent Algorithm:**

I. Start with random  $w^{t=0}$ 



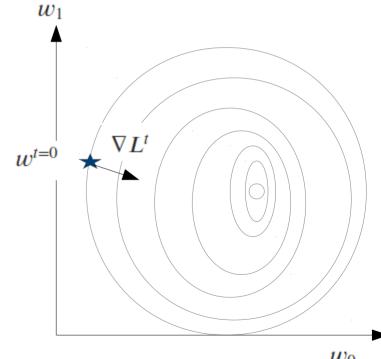


Goal: find w to minimize  $\arg \min \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

## **Gradient Descent Algorithm:**

- I. Start with random  $w^{t=0}$
- II. Compute gradient for all training samples

$$\nabla L^t = \sum_{i=0}^{|(X,y)|} \frac{dL(y_i, w^t x_i)}{dw^t}$$





Goal: find w to minimize  $\underset{w}{\operatorname{arg \, min}} \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

## **Gradient Descent Algorithm:**

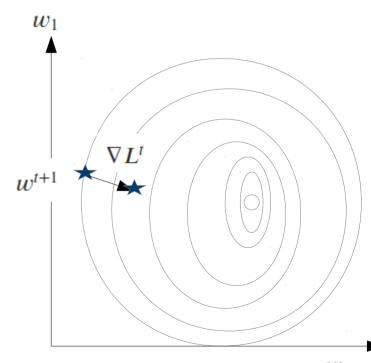
- I. Start with random  $w^{t=0}$
- II. Compute gradient for all training samples

$$\nabla L^t = \sum_{i=0}^{|(X,y)|} \frac{dL(y_i, w^t x_i)}{dw^t}$$

III. Update parameters

$$w^{t+1} = w^t + \lambda \nabla L^t$$

Step size or Learning rate
Usually quite small scalar like 0.001
ML Summer School -Janis Keuper





Goal: find w to minimize  $\underset{w}{\operatorname{arg \, min}} \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

## **Gradient Descent Algorithm:**

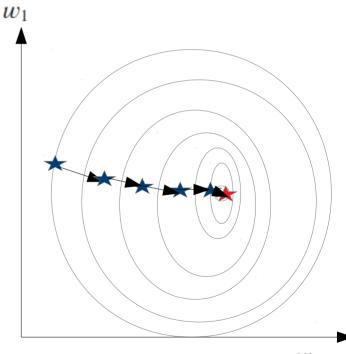
- I. Start with random  $w^{t=0}$
- II. Compute gradient for all training samples

$$\nabla L^t = \sum_{i=0}^{|(X,y)|} \frac{dL(y_i, w^t x_i)}{dw^t}$$

III. Update parameters

$$w^{t+1} = w^t + \lambda \nabla L^t$$

IV. Repeat II-III till convergence

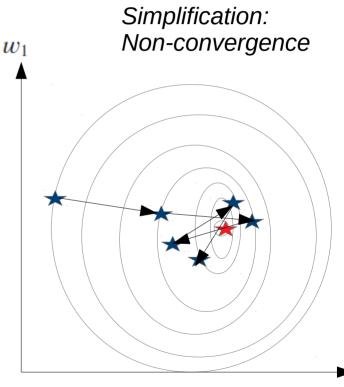




Goal: find w to minimize 
$$\underset{w}{\operatorname{arg \, min}} \sum_{i=0}^{N} L(y_i, w^T x_i)$$

#### Convergence

I. Theory: need to decrease  $\lambda$  to guarantee convergence to minimum

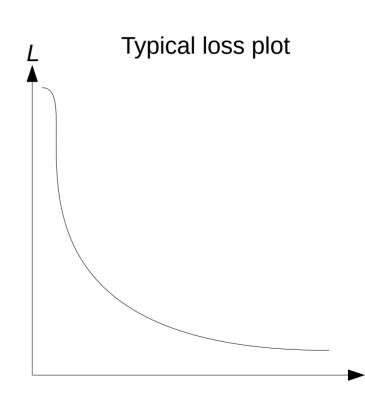




Goal: find w to minimize  $\underset{w}{\operatorname{arg \, min}} \sum_{i=0}^{N} L(y_i, w^T x_i)$ 

## Convergence

- I. Theory: need to decrease  $\lambda$  to guarantee convergence to minimum
- II. How to know when to stop?
  - I. Pre set number of iterations
  - II. Loss limit
  - III.Loss not changing



# **Multi Class Problems**



## What if we have more than two classes? → simple extension of our model

$$f(x) = y' = argmax(Wx)$$



- → One vector per class
- → Matrix vector Multiplication
- → returns vector with class-wise response
- → argmax selects maximum class label

## **Multi Class Problems**



What if we have more than two classes? → simple extension of our model

$$f(x) = y' = argmax(Wx)$$

## Optimization problem is almost the same

$$\arg\min_{w} \sum_{i=0}^{N} L(y_i, Wx_i)$$

Change Loss to SOFTMAX function to normalize sum aver all probabilities to one

$$L(y^{i}, y^{i'}) := \frac{e^{y^{i'}}}{\sum_{i}^{k} e^{y^{j'}}}$$

Use "one-hot" coding of y

# **Multi Class Problems**



## What if we have more than two classes? → simple extension of our model

$$f(x) = y' = argmax(Wx)$$

## Optimization problem is almost the same

$$\arg\min_{w} \sum_{i=0}^{N} L(y_i, Wx_i)$$

## Change Loss to SOFTMAX function to normalize sum aver all probabilities to one

$$L(y^{i}, y^{i'}) := \frac{e^{y^{i'}}}{\sum_{j}^{k} e^{y^{j'}}}$$

Use "one-hot" coding of y ◀

Y is now a vector with k (number of classes) entries and  $y^i$  is the kth class label

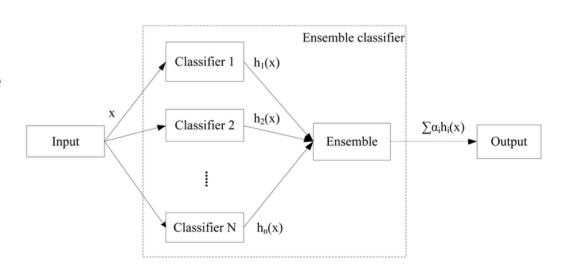


# **Discussion**



#### **Ensemble Learning**

- Very popular method based on ensemble learning
  - → many weak models decide together (by voting)
- Simple but powerful method
- Easy to implement and to parallelize
- Does not tend to overfit
- Build in Feature-Selection (next Lecture)

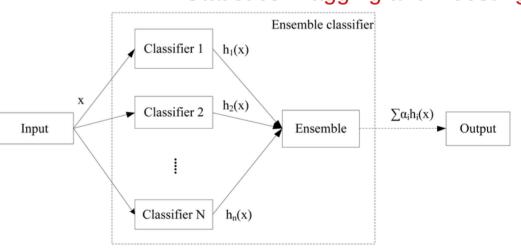




#### **Ensemble Learning**

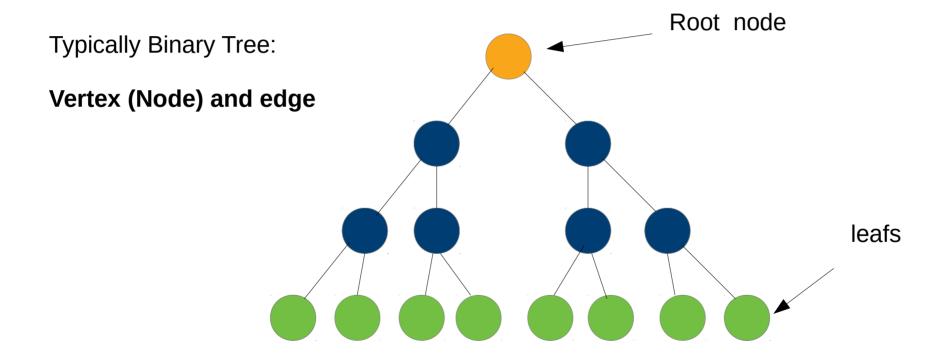
- Very popular method based on ensemble learning
  - → many weak models decide together (by voting)
- Simple but powerful method
- Approximating non-linear decision function by combination of piecewise linear functions
- Easy to implement and to parallelize
- Build in Feature-Selection (next Lecture)

Statistics: Bagging and Boosting





#### **Decision Trees: the base classifier for Random Forests**

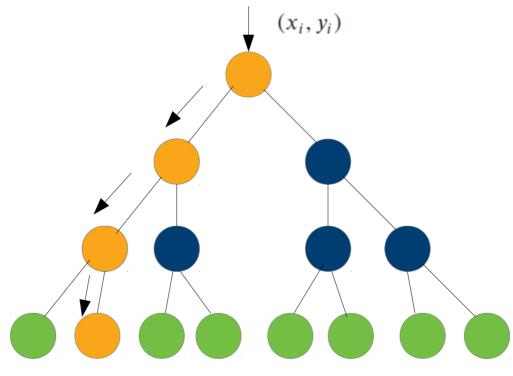




Goal: divide training data samples X at each node such that leafs have (mostly)

Pure class labels

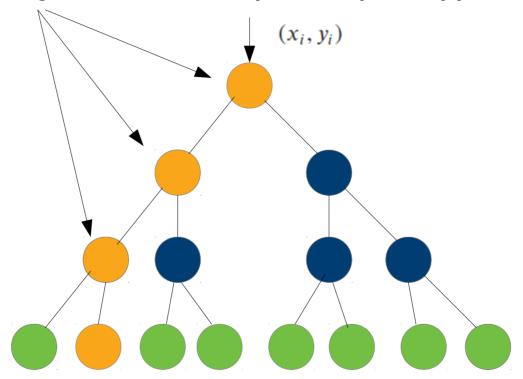
Classification: label y' is the label of the leaf node



$$\forall y_i, y_j \in L_1 : y_i = y_j$$



#### All we need is a splitting function that will produce (almost) pure class labels



$$\forall y_i, y_j \in L_1 : y_i = y_j$$



#### **Top Down Training:**

Assume k classes and data of dimension d

- Fill tree with ALL training samples from the root down
- In each node: compute probability for all class labels in node n

$$p_i := p(y = i) = \frac{|(x,i)| \in X_n}{|X_n|}$$

- Compute node purity based on class probabilities
- Split node along one dimension such that purity of children is increasing ← optimization



All we need is a splitting function that will produce (almost) pure class labels.

#### **Entropy (a way to measure impurity):**

$$Entropy = -\sum_{j} p_{j} \log_{2} p_{j}$$

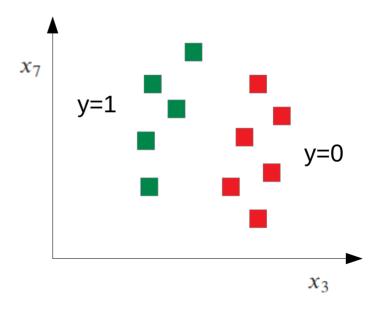
#### **Gini index:**

$$Gini = 1 - \sum_j p_j^2$$



#### **Split optimization** is a simple line search (Example):

I. Select a random subset of variables (feature dimensions) from the data X e.g.  $x_7$   $x_3$ 





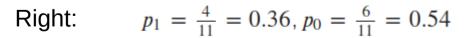
#### Split optimization is a simple line search (Example):

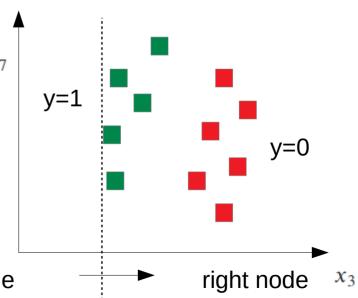
I. Select a random subset of variables (feature dimensions) from the data X e.g.  $x_7$   $x_3$ 

II. For each variable: find best split via line search e.g. for  $x_3$ 

Left:

Undefined (div by zero)





Left node

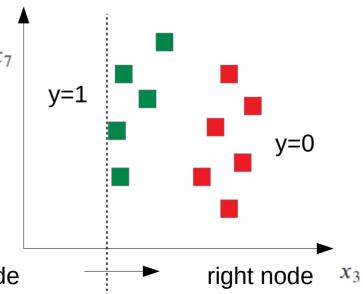


#### Split optimization is a simple line search (Example):

- I. Select a random subset of variables (feature dimensions) from the data X e.g.  $x_7$   $x_3$
- II. For each variable: find best split via line search e.g. for  $x_3$

Left: 
$$gini = 1 - (0^2 + 0^2) = 1$$

Right: 
$$gini = 1 - (0.36^2 + 0.54^2) = 0.57$$



Left node

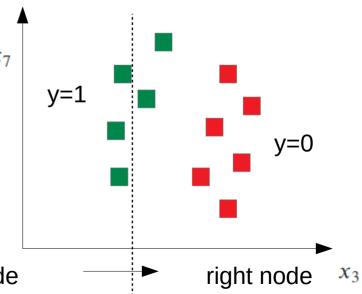


#### Split optimization is a simple line search (Example):

- I. Select a random subset of variables (feature dimensions) from the data X e.g.  $x_7$   $x_3$
- II. For each variable: find best split via line search e.g. for  $x_3$

Left: 
$$p_1 = \frac{3}{3} = 1, p_0 = \frac{0}{3} = 0$$

Right: 
$$p_1 = \frac{2}{8} = 0.25, p_0 = \frac{6}{8} = 0.75$$



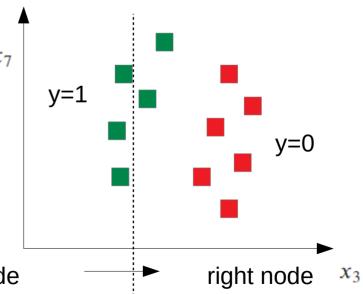


#### Split optimization is a simple line search (Example):

- I. Select a random subset of variables (feature dimensions) from the data X e.g.  $x_7$   $x_3$
- II. For each variable: find best split via line search e.g. for  $x_3$

Left: 
$$gini = 1 - (1^2 + 0^2) = 0$$

Right: 
$$gini = 1 - (0.25^2 + 0.75^2) = 0.375$$



Left node

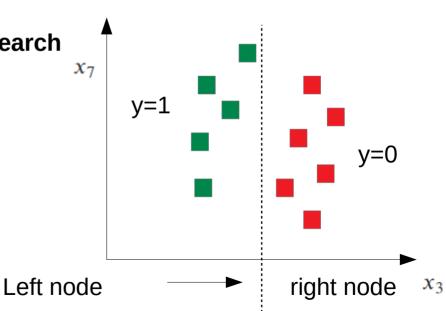


#### Split optimization is a simple line search (Example):

- I. Select a random subset of variables (feature dimensions) from the data X e.g.  $x_7$   $x_3$
- II. For each variable: find best split via line search e.g. for  $x_3$

Left: 
$$p_1 = \frac{5}{5} = 1, p_0 = \frac{0}{5} = 0$$

Right: 
$$p_1 = \frac{0}{6} = 0, p_0 = \frac{6}{6} = 1$$



ML Summer School -Janis Keuper

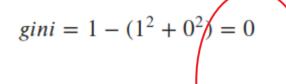


#### Split optimization is a simple line search (Example):

I. Select a random subset of variables (feature dimensions) from the data X e.g.  $x_7$   $x_3$ 

II. For each variable: find best split via line search e.g. for  $x_3$ 

Left:

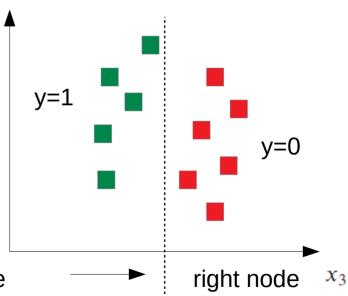


Right:

$$gini = 1 - (0^2 + 1^2) = 0$$

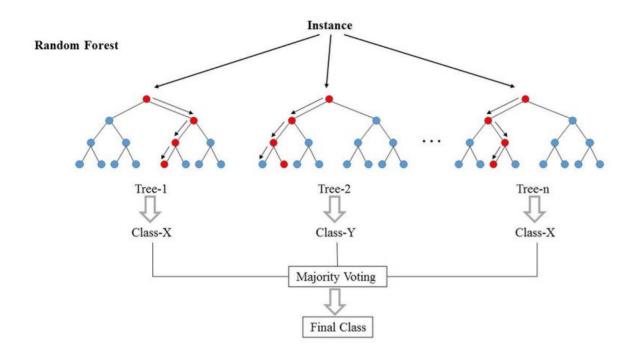
**Perfect split!** 

Left node



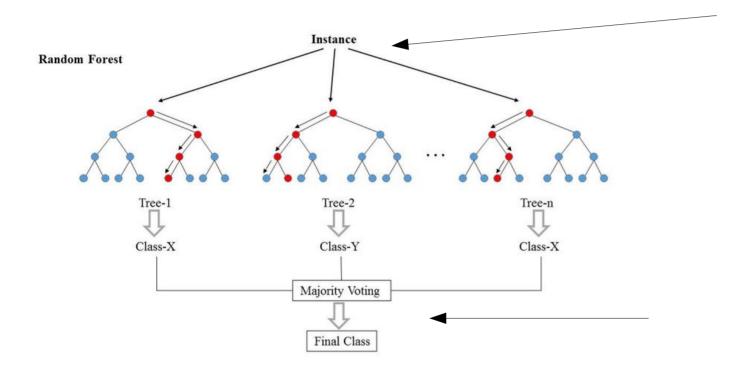


#### **Ensemble Learning: A Forest of Trees**





#### **Ensemble Learning: A Forest of Trees**



Split Training data into random subsets

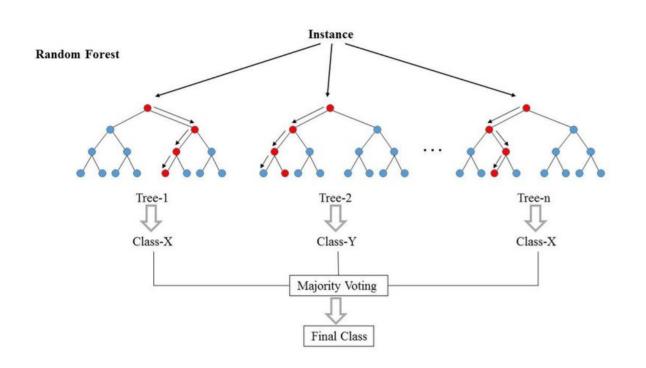
→ Bootstrap

**Combine Models** 

→ Bagging



#### **Ensemble Learning: A Forest of Trees**

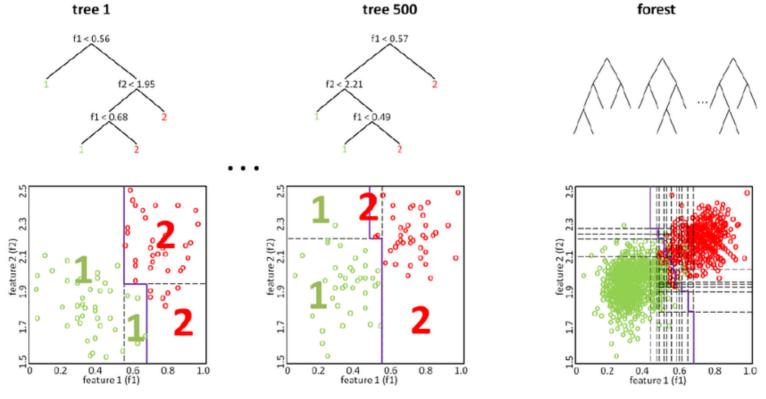


#### Parameters:

- #trees
- Portion of data per tree
- #vars per split
- Stopping
  - max depth
  - min samples per node



#### **Classification example**



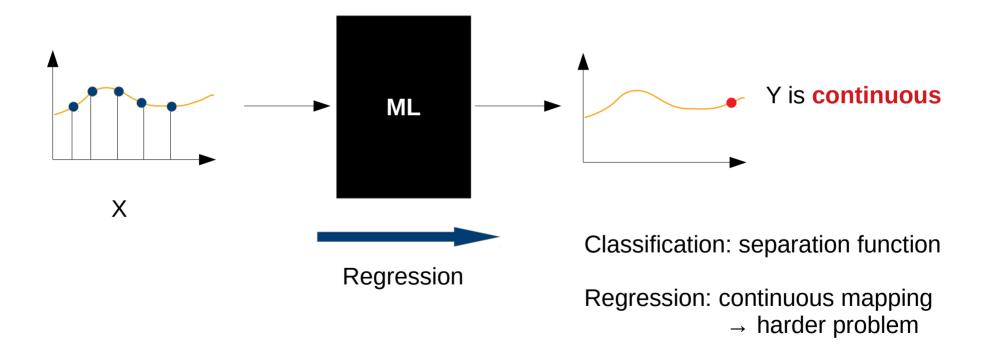


# **Discussion**

# Regression

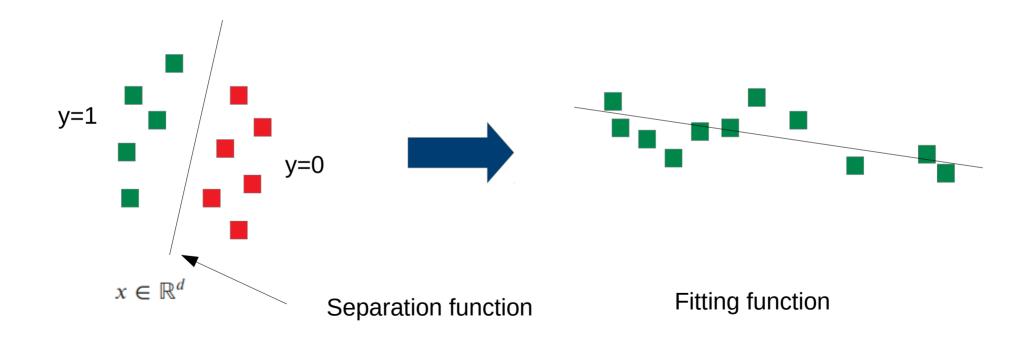


#### Recall:





How do we have to change out linear classifier to predict continuous values?

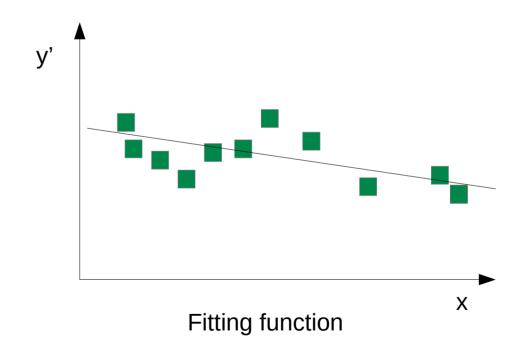




How do we have to change out linear classifier to predict continuous values?

#### Still can use the same framework

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$





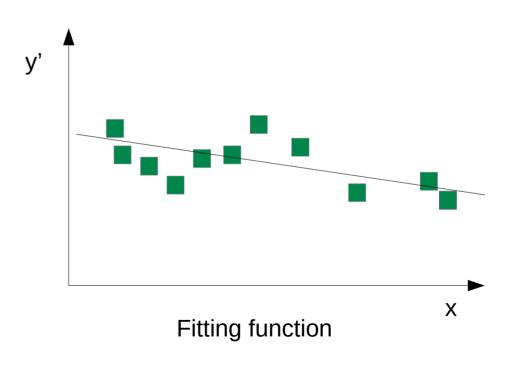
How do we have to change out linear classifier to predict continuous values?

#### Still can use the same framework

$$f(x) = y' = w^T x = \sum_{j=0}^d w_j x_j$$

# Simply need new loss function in the optimization

$$\underset{w}{\arg\min} \sum_{i=0}^{N} L(y_i, w^T x_i)$$





#### **Loss functions for regression:**

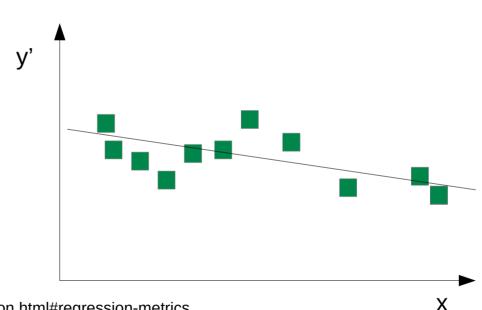
$$\underset{w}{\arg\min} \sum_{i=0}^{N} L(y_i, w^T x_i)$$

#### As simple as least squares error

$$L_{LSE}(y, y') := ||y - y'||^2$$

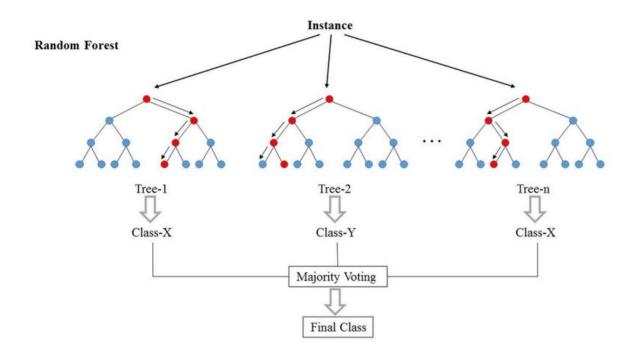
#### Many other error measures possible

- L1 (Histogram intersection)
- ...
  - → See https://scikit-learn.org/stable/modules/model\_evaluation.html#regression-metrics



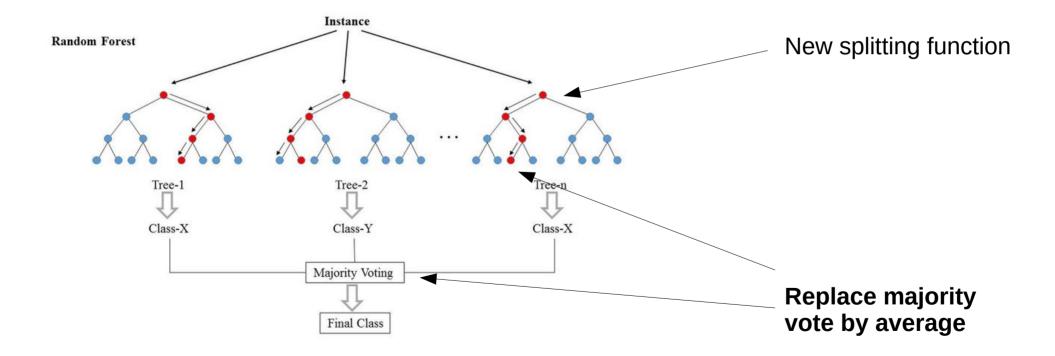


#### **Recall:**





#### **Recall:**





#### **Splitting functions for regression:**

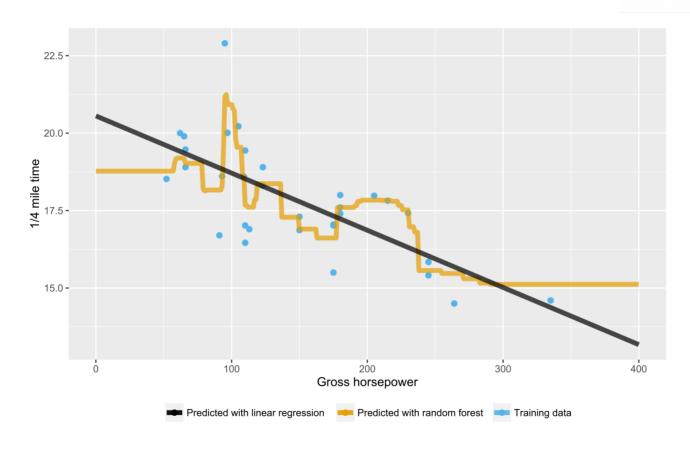
Goal: reduce "data spread" in node

→ use simple statistical measure like "mean square error"

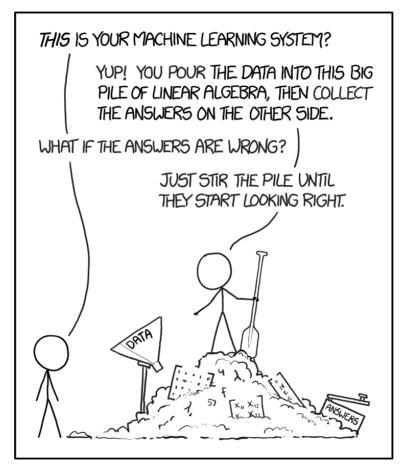
$$MSE : \sum_{(x_i, y_i) \in X_n} \|\mu_y - y_i\|^2$$



#### **Example**







https://xkcd.com/1838/